

A Review on Knowledge Graphs for Healthcare: Resources, Applications, and Promises

Hejie Cui¹, Jiaying Lu¹, Shiyu Wang^{2†}, Ran Xu^{1†}, Wenjing Ma^{3†},
Shaojun Yu^{1†}, Yue Yu^{4†}, Xuan Kan^{1†}, Chen Ling¹, Tianfan Fu⁵,
Liang Zhao¹, Joyce Ho¹, Fei Wang^{6,7}, Carl Yang^{1*}

¹Department of Computer Science, Emory University, Atlanta, GA, USA.

²Department of Biostatistics, Emory University, Atlanta, GA, USA.

³Department of Biostatistics, University of Michigan, Michigan, MI, USA.

⁴School of CSE, Georgia Institute of Technology, Atlanta, GA, USA.

⁵Department of Computer Science, Rensselaer Polytechnic Institute, Troy, NY, USA.

⁶Weill Cornell Medicine, New York, NY, USA.

⁷Department of Population Health Science, Cornell University, New York, NY, USA.

*Corresponding author(s). E-mail(s): j.carlyang@emory.edu;

†These authors contributed equally to this work.

Abstract

Healthcare knowledge graphs (HKGs) are valuable tools for organizing biomedical concepts and their relationships with interpretable structures. The recent advent of large language models (LLMs) has paved the way for building more comprehensive and accurate HKGs. This, in turn, can improve the reliability of generated content and enable better evaluation of LLMs. However, the challenges of HKGs such as regarding data heterogeneity and limited coverage are not fully understood, highlighting the need for detailed reviews. This work provides the first comprehensive review of HKGs. It summarizes the pipeline and key techniques for HKG construction, as well as the common utilization approaches, i.e., model-free and model-based. The existing HKG resources are also organized based on the data types they capture and application domains they cover, along with relevant statistical information.* At the application level, we delve into the successful integration of HKGs across various health domains, ranging from fine-grained basic science research to high-level clinical decision support and public health. Lastly, the paper highlights the opportunities for HKGs in the era of LLMs. This work aims to serve as a valuable resource for understanding the potential and opportunities of HKG in health research.

*Resource available at <https://github.com/lujiaying/Awesome-HealthCare-KnowledgeBase>

Keywords: healthcare, knowledge graph, review, interpretability, multimodality

1 Introduction

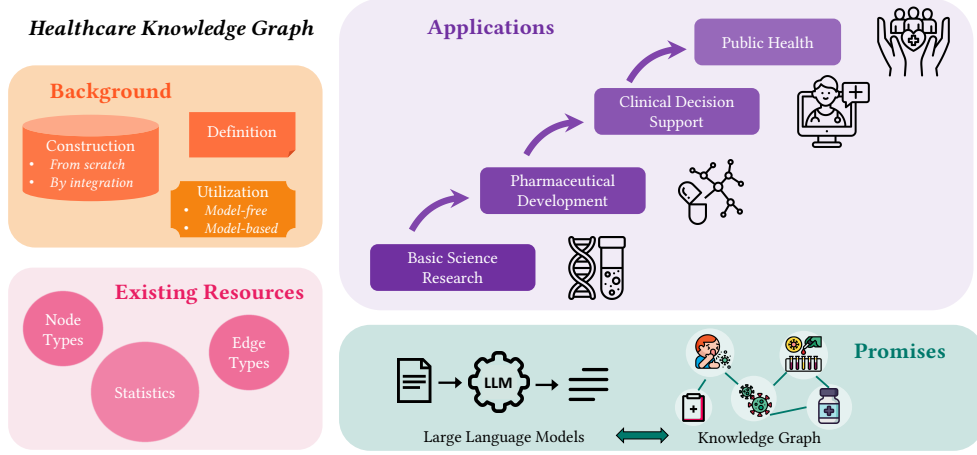


Fig. 1 The content overview of this review on healthcare knowledge graph.

A knowledge graph (KG) is a data structure that captures the relationships between different entities and their attributes [1, 2]. KG models and integrates data from various sources, including structured and unstructured data, and has been studied to support a wide range of applications such as search engines [3], recommendation systems [4, 5], and question answering [6–9]. Healthcare Knowledge Graph (HKG) facilitates an interpretable representation of medical concepts, e.g., drugs and disease, as well as the relations among those medical concepts. This data structure enables the connection of contexts and enhances clinical research and decision-making [10, 11].

On the data side, HKG is usually built on complex medical systems such as electronic health records, medical literature, clinical guidelines, and patient-generated data [12, 13]. However, these data resources are often heterogeneous and distributed, which makes it challenging to integrate and analyze them effectively [14]. This data heterogeneity can also lead to incomplete or inconsistent data representations, limiting their usefulness for downstream healthcare tasks [15]. Additionally, the current use of domain-specific knowledge graphs may result in limited coverage and granularity of the knowledge captured across different levels. This hinders identifying correlations and relationships between medical concepts from multiple domains. These challenges highlight the need for continued research on HKGs to realize their full potential.

On the modeling side, the construction of HKGs can be done either from scratch or by integrating existing dataset resources. Many crucial steps, such as entity and

relation extraction, can be optimized with natural language processing tools and algorithms. Recently, there have been significant advancements in general domain knowledge extraction, thanks to pre-trained large language models (LLMs) such as BERT [16], GPT Series [17], and others. These models revolutionize the field and make it possible to integrate heterogeneous medical data from various sources effectively. The use of pre-trained models has also led to the development of more accurate and comprehensive medical ontologies and taxonomies [18–22]. This allows for the evaluation of generated contents from LLMs and reduces LLM hallucination.

A comprehensive healthcare knowledge graph has the potential to contribute to health research across various levels [10, 23, 24]. At the micro-scientific level, HKGs can help researchers identify new phenotypic and genotypic correlations and understand the underlying mechanisms of disease [25], leading to more targeted and effective treatments [11, 26]. At the clinical care level, HKGs can be used to develop clinical decision support systems that provide clinicians with relevant information, improving clinical workflows and patient outcomes [27, 28]. Therefore, a thorough review of existing literature on HKGs becomes an essential roadmap and invaluable resource to drive transformative advancements in the field.

This review is the first comprehensive overview of HKGs, covering contents shown in Figure 1. Specifically, in Section 2, we delve into the construction pipelines of HKGs, including both building from scratch and integration approaches, and highlight the key techniques employed in HKG construction. Additionally, we explore two standard utilization methods of HKGs, namely model-free and model-based approaches. In Section 3, we compile a comprehensive summary of existing HKG resources with their scopes and applications. Furthermore, Section 4 investigates the literature on mainstream health applications, offering an in-depth overview of the diverse use cases of HKGs in healthcare. Finally, we address promising research opportunities in the era of LLMs in Section 5.

2 Backgrounds

2.1 Knowledge Graphs for Healthcare (HKG)

Definition. A healthcare knowledge graph (HKG) is a domain-specific knowledge graph designed to capture medical concepts such as drugs, diseases, genes, phenotypes, and so on, and their relationships in a structured and semantic way.

Terminology. Our focus is on healthcare knowledge graphs (HKGs), which are structured, semantic representations of medical concepts and relationships. We also include ontologies and knowledge bases, which are commonly used in constructing HKGs. Ontologies define a set of concepts and categories in a domain, as well as the relationships between them, while knowledge bases store factual information about entities and their attributes. By covering these categories of terminology, we provide a comprehensive overview of the different types of resources available for organizing and representing medical knowledge in a structured and semantically rich manner.




Fig. 2 Detailed taxonomy of the background section of healthcare knowledge graphs.

2.2 HKG Construction

Healthcare knowledge graphs can be constructed from scratch or through the integration of existing data resources.

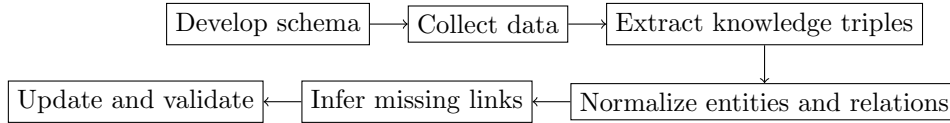


Fig. 3 The pipeline of constructing HKGs from scratch.

Constructing HKGs from Scratch. A multi-step pipeline, as in Figure 3, is used to construct HKGs from scratch.

1. The first step is to identify the scope and objectives. In most cases, researchers develop a schema [29, 30] or use existing schemas [29, 31–33] to serve as the formal and explicit specification of a domain, thus ensuring consistent, coherent, and aligned domain knowledge. Unlike the general domain KG, utilizing schemas is a common practice in HKG construction.
2. Secondly, researchers gather data from various sources, including medical literature, clinical trials, and patient-generated data. It is essential to ensure the quality and consistency of the data and to remove identifiable information for patient privacy.
3. The third step is to extract and transform the data into a structured format. This step involves identifying medical entities and creating relationships between them via specialized biomedical Natural Language Processing (NLP) tools [34–36].
4. Next, researchers map the entities and relationships to the chosen ontologies with the help of thesauruses [37] or terminologies [38, 39]. This step ensures that the knowledge graph is interoperable with other healthcare systems and facilitates data integration.
5. Until now, an initial KG has been built. The next step is to populate the KG to infer missing links between entities. This inference can be done using graph databases [40] or link prediction models [41, 42].
6. The final step is to continuously update and validate the KG to ensure accuracy and relevance. This step involves incorporating new data and knowledge, refining the schema, and evaluating the quality of the KG.

Constructing HKGs by Integration. Considering significant efforts have been paid to construct and curate HKGs from scratch, it is promising to integrate these data resources to avoid repetitive work. Healthcare KG integration (also called Healthcare KG fusion) refers to the processing of merging two or more HKGs into a single, more comprehensive graph [43–45]. The integration process is challenging because different HKGs may use different terminologies, schemas, or data formats. To address these challenges, researchers have developed various techniques and algorithms for knowledge graph fusion, including ontology matching [46, 47], schema alignment [48, 49], entity resolution [50, 51], and conflicts resolution [52]. These methods aim to identify and reconcile the differences between KGs.

Techniques for HKG Constructions. Traditionally, each step of HKG construction involves one specially designated model. For instance, Hidden Markov Models and Recurrent Neural Networks are widely used for healthcare named entity recognition, relation extraction, and other sequence tagging tasks, while Translational Models and Graph Neural Networks are used for HKG completion and conflicts resolution tasks. Recently, large language models (LLMs) have shown great utility to serve as a uniform tool for constructing KGs [53]. Several key steps of constructing KGs, such as named entity recognition [54–57], relation extraction [58–60], entity linking [61–63], and KG completion [64–66], have been successfully tackled by these large foundation models. Early explorations of construction HKG with large foundation models show that healthcare entity normalization [67, 68], healthcare entity recognition [69, 70], healthcare entity linking [71], and healthcare knowledge fusion [72] can also be performed, without extensive training on expensive healthcare annotated corpus. On the other hand, researchers start to construct KGs under the open-world assumption [42, 73–76], thus getting rid of the dependency on pre-defined schemas and exhaustive entity&relation normalization. Although open-world KGs greatly increase the coverage, ensuring the quality of extracted knowledge is still an open research challenge, especially for explainable and trustworthy HKGs.

2.3 HKG Utilization

Model-free Utilization. Various query languages can be used for KGs, such as SPARQL, Cypher, and GraphQL [40]. These query languages allow users to query healthcare KGs using a standardized syntax, thus enabling users to retrieve, manipulate, and analyze data in a structured and consistent way. More complex applications can be further supported by graph queries. For instance, automatic healthcare question answering can be tackled by Natural Language Question-to-Query (NLQ2Query) approach [90], where natural language questions are first translated into executable graph queries and then answered by the query responses. HKGs can also be utilized as an up-to-date and trustworthy augmentation to large language models (LLMs) for many applications. Some pioneering studies [91–94] show that retrieved knowledge triples can improve the reliability of LLMs in various knowledge-intensive tasks, by addressing the nonsensical or unfaithful generation. Moreover, KGs can be a useful tool for fact-checking [95–97] as they provide a structured representation of information that can be used to quickly and efficiently verify the accuracy of claims. Researchers

Table 1 Resource of existing healthcare knowledge graph (HKG).

Name	Node Types	Edge Types	Statistic	Application
HetioNet [77]	11 (e.g., drug, disease)	24 (e.g., drug-disease)	#N: 47.0 K, #E: 2.3 M	Medicinal Chemistry
DrKG [78]	13 (e.g., disease, gene)	107 (e.g., disease-gene)	#N: 97 K, #E: 5.8 M	Medicinal Chemistry
PrimeKG [11]	10 (e.g., phenotypes)	30 (e.g., disease-phenotype)	#N: 129.4 K, #E: 8.1 M	Medicinal Chemistry
Gene Ontology ^a [31]	3 (e.g., biological process)	4 (e.g., partOf)	#N: 43 K, #E: 7544.6K	Bioinformatics
KEGG ^b [?]]	16 (e.g., pathway)	4 (e.g., partOf)	#N: 48.5 M, #E: unknown	Bioinformatics
STRING ^c [?]]	1 (e.g., protein)	4 (e.g., interactions)	#N: 67.6 M, #E: 20 B	Bioinformatics
Cell Ontology ^d [?]]	1 (i.e., cell type)	2 (e.g., subclassOf)	#N: 2.7 K, #E: 15.9 K	Bioinformatics
GEFA [79]	510 (e.g., kinases)	2 (e.g., drug-drug)	#N: 0.5 K, #E: 30.1 K	Drug Development
Reaction [80]	2 (e.g., reactant & normal)	19 (e.g., reaction paths)	#N: 2192.7 K, #E: 932.2 K	Drug Development
ASICS [81]	2 (e.g., reactant & product)	1 (e.g., reactions)	#N: 1674.9 K, #E: 923.8 K	Drug Development
Hetionet [81]	11 (e.g., biological process)	24 (e.g., disease-associates-gene)	#N: 47.0 K, #E: 2250.2 K	Drug Development
LBD-COVID [82]	1 (i.e., concept)	1 (i.e., SemMedDB relation)	#N: 131.4 K, #E: 1016.1 K	Drug Development
GP-KG [83]	7 (e.g., drug)	9 (e.g., disease-gene)	#N: 61.1 K, #E: 1246.7 K	Drug Development
DRKF [84]	4 (e.g., drug)	43 (e.g., drug-disease)	#N: 12.5 K, #E: 165.9 K	Drug Development
DDKG [85]	2 (i.e., drug & disease)	1 (e.g., drug-disease)	#N: 551, #E: 2.7 K	Drug Development
Disease Ontology ^e [32]	1 (i.e., disease)	2 (e.g., subclassOf)	#N: 11.2 K, #E: 8.8 K	Clinical Decision Support
DrugBank [86]	4 (e.g., drug, pathway)	4 (e.g., drug-target)	#N: 7.4 K, #E: 366.0 K	Clinical Decision Support
KnowLife [87]	6 (e.g., genes)	14 (e.g., gene-diseases)	#N: 2.9 M, #E: 11.4 M	Clinical Decision Support
PharmKG [88]	3 (e.g., diseases)	3 (e.g., chemical-diseases)	#N: 7601, #E: 500958	Clinical Decision Support
ROBOKOP ^f [89]	54 (e.g., genes, drugs)	1064 (e.g., biolink, CHEBI)	#N: 8.6M, #E: 130.4 M	Clinical Decision Support
iBKH ^g [44]	11 (e.g., anatomy, disease)	18 (e.g., anatomy-gene)	#N: 2.4 M, #E: 48.2 M	Clinical Decision Support

^a<http://geneontology.org/>

^b<https://www.genome.jp/kegg/>

^c<https://string-db.org/>

^d<https://www.ebi.ac.uk/ols4/ontologies/cl>

^e<https://disease-ontology.org/>

^f<https://robokop.renci.org/>

^g<https://github.com/wcm-wanglab/iBKH>

have explored the utility of HKGs in identifying ingredient substitutions of food [98], COVID-19 fact-checking [99], etc.

Model-based Utilization. Utilizing HKGs in complex reasoning tasks often involves utilizing machine learning models. HKG embeddings [100, 101] have shown great potential to tackle these tasks. In particular, HKG embedding models are a class of machine learning models that aim to learn low-dimensional vector representations, or embeddings, of the entities and relations in a knowledge graph. After obtaining HKG embeddings, they can be plugged into any kind of deep neural network and further fine-tuned toward downstream objectives. On the other hand, symbolic logic models represent another prominent approach for KG reasoning due to their interpretability. More specifically, symbolic reasoning models first mine logical rules from existing knowledge by inductive logic programming [102], association rule mining [103], or Markov logic networks [104]. These mined rules are used to infer new facts, make logical deductions and answer complex queries. Recently, researchers start to explore combining logical rules into KG embedding to further improve the generalization and performance of HKG reasoning [105, 106].

3 Resources

In this section, we compile a detailed resource overview of existing healthcare knowledge graphs to assist researchers and healthcare professionals in constructing and utilizing HKGs, organized in Table 1. We present key attribute information, including HKG name, node types, edge types, statistics, and their applications. Details and external links can be referred to at our maintained public repository ¹.

4 Applications

4.1 Basic Science Research

Several previous biological terms can also be considered knowledge graphs, such as ontology (gene ontology, cell ontology, disease ontology), network (gene regulatory network), etc. We use the original biological terms as they are more popular according to historical reasons.

4.1.1 Medicinal Chemistry

Topics related to medicinal chemistry involve drug-drug interactions (DDIs) and drug-target interactions (DTIs), which will be discussed in this section.

Drug-drug interactions (DDIs) refer to changes in the actions, or side effects, of drugs when they are taken at the same time or successively [107]. In general, DDIs are a significant contributor to life-threatening adverse events [101], and their identification is one of the key tasks in public health and drug development. The existence of diverse datasets on drug-drug interactions (DDIs) and biomedical KGs has enabled the development of machine-learning models that can accurately predict DDIs. Yu et al. [100] develop SumGNN, a model that includes a subgraph extraction module to

¹<https://github.com/lujiaying/Awesome-HealthCare-KnowledgeBase>




Fig. 4 Detailed taxonomy of the application section of healthcare knowledge graphs.

efficiently extract relevant subgraphs from a KG, a self-attention-based summarization scheme to generate reasoning paths within the subgraph, and a multichannel module for integrating knowledge and data, resulting in significantly improved predictions of multi-typed DDIs. Su et al. [101] propose DDKG, an attention-based KG representation learning framework that involves an encoder-decoder layer to learn the initial embeddings of drug nodes from their attributes in the KG. Karim et al. [108] compare various techniques for generating KG embeddings with different settings and conclude that a combined convolutional neural network and LSTM network yields the highest accuracy when predicting drug-drug interactions (DDIs). Dai et al. [109] propose a new KG embedding framework by introducing adversarial autoencoders based on Wasserstein distances and Gumbel-Softmax relaxation for DDI tasks. Lin et al. [110] develop KGNN that resolves the DDI prediction by capturing drug and its potential neighborhoods by mining their associated relations in KG.

Drug-target interactions (DTIs) is just as important as DDIs [111]. Machine learning models can leverage knowledge graphs constructed from various types of interactions, such as drug-drug, drug-disease, protein-disease, and protein-protein interactions, to aid in predicting DTIs. For instance, Li et al. [112] utilize the KG transfer probability matrix to redefine the drug-drug and target-target similarity matrix, thus constructing the final graph adjacent matrix to learn node representations by VGAE and augmenting them by utilizing dual Wasserstein Generative Adversarial Network with gradient penalty. Zhang et al. [113] propose a new hybrid method for DTI prediction by first constructing DTI-related KGs and then employing graph representation learning model to obtain feature vectors of the KG. Wang et al. [114]

construct a knowledge graph of 29,607 positive drug-target pairs by DistMult embedding strategy, and propose a Conv-Conv module to extract features of drug-target pairs. Ye et al. [115] learn a low-dimensional representation for various entities in the KG, and then integrate the multimodal information via neural factorization machine.

4.1.2 Bioinformatics Research

In bioinformatics settings, a knowledge graph is a graph-based representation where nodes are biomedical entities (such as mutations, genes, proteins, metabolites, diseases, and biological pathways), and edges are their relationships (such as associations, interactions, and regulations). Through the integration, researchers can gain a more comprehensive understanding of complex biological processes and diseases.

Multi-Omics Applications: In recent years, the field of multi-omics data analysis has become increasingly important for understanding complex biological systems. With the advancement of high-throughput technologies, researchers can generate large-scale and high-dimensional data from different omics fields, such as genomics, transcriptomics, proteomics, metabolomics, and epigenomics. More KG applications based on multi-omics data integration have emerged, aiming to provide new research methods to uncover the complex relationships between different omics layers and reveal biological systems’ underlying mechanisms.

Knowledge graphs have been used to identify disease-associated mutations, genes, proteins, and metabolites by integrating multi-omics data with existing biological knowledge. This approach has led to the discovery of novel biomarkers and therapeutic targets for various diseases and the interpretation of the functional effects of genetic elements [116]. Quan et al. built a comprehensive multi-relational knowledge graph called AIMedGraph, providing an interpretation of the impact of genetic variants on disease or treatment [117]. They curated detailed information about diseases, drugs, genetic variants, and the impact of genetic variations on disease development and drug treatment from multiple data resources. The entities integrated into AIMedGraph are connected by evidence-based relations and form a comprehensive gene-variant-disease-drug-trial-reference knowledge network. AIMedGraph uses the Adamic-Adar algorithm to predict new relations between entities based on shared neighbors and their proximity in the knowledge network. GenomicsKG is a knowledge graph to analyze and visualize multi-omics data. GenomicsKG can be used to improve drug development based on clinical genomics correlations and personalized drug customization in the extended version based on interactive relationships. It also provides multi-dimensional visualization, linked functional knowledge graphs, and reporting for clinical genomics.

Single-Cell Analysis: Cells are fundamental and essential units of living organisms. With high-throughput sequencing technologies advancing to measure genomic profiles in a single-cell resolution, cell functions (inside cells) and cell-cell interactions (between cells) are revealed [118]. When diving into the functions of cells, gene regulatory mechanisms can be a critical factor. Gene regulatory mechanisms control the expression of genes, which can affect cell differentiation, response to stimuli, disease progression, etc. It can be visualized as gene regulatory networks (GRNs) which depict the interactions between genes and their regulators. Traditional approaches

to constructing the GRN are based on gene knockdown or knockout experiments, which can be time-consuming, labor-intensive, and costly. In addition, these experiments are limited to only one or a few genes which neglects gene-gene interactions. In contrast, single-cell sequencing data could provide whole-genome scale measurements in each individual cell with comprehensive information such as gene expression, transcription factors (TF) binding sites, DNA methylation, epigenetic modifications, etc. By mining these publicly available data, it is possible to reveal the underlying and universal GRN biomedical researchers to understand biological processes better. For example, GRNdb provides detailed regulon and TF-target pairs information from different human and mouse tissues under different conditions by analyzing existing sequencing data [119]. GenomicKB integrates existing datasets and genome annotations and formulates the data into a knowledge graph to emphasize the relationships among genomic entries [120]. We can foresee that with more data being generated and collected, one generalized GRN (gene regulatory knowledge graph) along with cell-type-specific GRNs can be reconstructed and used to help reveal the underlying mechanisms of gene expression, biological processes, and disease progression.

Besides cell-type-specific GRN, single-cell sequencing data can also be used to infer cell-cell interactions, which play a critical role in understanding cell cycles, cell fate decisions, tissue development, immune response, etc. Among several approaches through which cells can interact, cell-cell communication or cell signaling is of the most interest [121] as the physical distance between cells does not limit it. Cells can communicate with each other by sending signaling molecules called ligands and receiving through receptors located on cell surfaces. Several methods have been developed to implicitly measure cell-cell communication by modeling and scoring ligand-receptor interactions (LRIs) from single-cell sequencing data [122, 123]. However, these methods do not utilize the spatial information of cells. Recent advancements in spatial sequencing techniques provide not only genomic measurements but also complementary information on the spatial coordinates of each cell, where the spatial coordinates indicate the probabilities of cells communicating with each other as cells with closer distances tend to interact more. A recent benchmark study on evaluating cell-cell communication inference methods based on spatial distances [124] while another benchmark study also considers other biological factors such as cytokine activities and receptor protein abundance [125]. With the availability of spatial transcriptomics data, many methods have been developed to leverage colocalization information to infer LRIs. Although many LRIs databases [122, 126] have been constructed and applied to infer cell-cell communication, only recently, SpaTalk [127] integrates CellTalkDB, Kyoto Encyclopedia of Genes and Genomes (KEGG) pathways, Reactome and TFs from AnimalTFDB to construct a ligand-receptor-target knowledge graph to help improve the inference of cell-cell communication. The success of SpaTalk provides an unprecedented view of integrating existing networks into a knowledge graph and shows that correctly utilizing the knowledge graph could provide helpful information in cell-cell communication tasks.

4.2 Pharmaceutical Research Development

4.2.1 Drug Development

Drug development is the process of identifying novel chemical compounds that can effectively treat or alleviate human diseases. Before the drug can be designated as a final product for clinical use, several critical steps need to be undertaken from the initial target identification, chemical synthesis and clinical trials. The whole process typically spans over a decade and involves expenses of approximately one billion dollars [128], yet it is characterized by a low success rate for clinical approval [129].

With advancements in chemistry and the availability of vast chemical libraries, it is possible to reduce the costs of drug development by extracting and integrating valuable insights from existing data and using computer algorithms to accelerate the drug design process [130–136]. Despite the growing trend of computer-assisted drug discovery, a key question remains regarding how to effectively integrate data and extract valuable insights from the vast chemical dataset.

To approach this question, knowledge graphs (KGs) have been employed for drug discovery due to their various advantages [131]: (1) In contrast to traditional methods that capture only one type of relationship, KGs are capable of providing heterogeneous information that includes diverse entities (*e.g.*, scaffolds, proteins, and genes); (2) In addition, KGs are capable of handling multiple types of relationships between various types of entities, such as drug-target pairs; and (3) KGs can provide unstructured semantic relationships between entities. In such graphs, entities are represented as nodes while their relationships are represented as edges, by which complex relations in biochemical systems can be easily handled.

In general, the field of drug development encompasses two main areas: *drug design* and *drug repurposing* (also known as drug repositioning). Drug design is the process of creating novel and diverse drug molecules with desirable pharmaceutical properties [137–139], whereas drug repurposing identifies new uses for existing approved drugs that were originally developed for a different indication [140, 141].

Drug Design. KGs are widely employed in drug design, particularly in generating novel molecules that are promising drug candidates for various diseases [79, 80]. Ranjan et al. [79] utilize Gated Graph Neural Network (GGNN) to generate novel molecules that target the coronavirus (*i.e.*, SARS-CoV-2) [142] and integrate KGs into their approach to reduce the search space. Specifically, KGs were leveraged to discard non-binding molecules before inputting them into the Early Fusion model, thus optimizing the efficiency of the drug design process. In addition to employing deep learning for direct structure design, KGs are also utilized in the analysis of chemical synthesis. Quantitative estimation of molecular synthetic accessibility is critical in prioritizing the molecules generated from generative models. For instance, Li et al. [80] utilize reaction KGs to construct classification models for compound synthetic accessibility. By leveraging KGs that capture information about reactions, including reaction types, substrates, and reaction conditions, they can train machine learning models that could predict the synthetic accessibility of compounds. Jeong et al. [81] introduce an intelligent system that integrates generative exploration and exploitation of reaction knowledge base to support synthetic path design.

Drug Repurposing. Compared to drug design, KGs are more commonly utilized to expedite the drug re-purposing process [43, 82–85, 143–145]. Many applications on drug re-purposing that utilize KGs are primarily focused on link prediction tasks [144]. To re-purpose promising drug candidates for new indications, many methods employ predictive models that focus on predicting drug-treats-disease relationships within pharmacological knowledge graphs KGs. Himmelstein et al. [43] use a degree-normalized pathway model on the hetionet KG, which includes genes, diseases, tissues, pathophysiologicals, and multimodal edges, to identify potentially repurposable drugs for epilepsy. Xu et al. [145] develop a multi-path random walk model on a network that incorporates gene-phenotype associations, protein-protein interactions, and phenotypic similarities for training and prediction purposes. Zhang et al. [82] introduce an integrative and literature-based discovery model for identifying potential drug candidates from COVID-19-focused research literature, including PubMed and other relevant sources. Gao et al. [83] construct a knowledge graph (KG) by integrating multiple genotypic and phenotypic databases. They then learn low-dimensional representations of the KG and utilize these representations to infer new drug-disease interactions, providing insights into potential drug repurposing opportunities. Zhang and Che [84] introduce a model for drug re-purposing in Parkinson’s disease that leverages a local medical knowledge base incorporating accurate knowledge along with medical literature containing novel information. Ghorbanali et al. [85] present the DrugRep-KG method, which utilizes a KG embedding approach for representing drugs and diseases in the process of drug repurposing.

4.2.2 Clinical Trial

The major goal of clinical trials is to assess the safety and effectiveness of drug molecules on human bodies. A novel drug molecule needs to pass three phases of clinical trials before it is approved by the Food and Drug Administration (FDA) and enters the drug market. The whole process is prohibitively time-consuming and expensive, costing 7-11 years and two billion dollars on average [146].

Clinical Trial Optimization targets identifying eligible patients for clinical trials based on their medical history and health conditions [147, 148]. Recently, with massive electronic health records (EHR) data and trial eligibility criteria (EC), data-driven methods have been studied to automatically assign appropriate patients for clinical trials [149–151]. However, it is often hard to fully capture and represent the complex knowledge present in unstructured ECs and EHR data, as ECs may only provide general disease concepts. In contrast, patient EHR data contain more specific medical codes to represent patient conditions. To better capture the interactions among different medical concepts from EHR records and ECs, Gao et al. [152] enhance patient records with hierarchical taxonomies to align medical concepts of varying granularity between EHR codes and ECs. Besides, Fu et al. [153] leverage additional knowledge-embedding modules along with drug pharmacokinetic and historical trial data to improve the patient trial optimization process, and Wang et al. [154] leverage the knowledge graphs to learn static trial embedding and further designed meta-learning module to generalize well over the imbalanced clinical trial distribution.

4.3 Clinical Decision Support

Nowadays, abundant Electronic Health Record (EHR) data enables better computational models for accurate diagnoses and treatments. EHR contains essential patient information such as disease diagnoses, prescribed medications, and test results. Due to this valuable information, EHRs are extensively utilized to identify patterns in patient health and assist healthcare providers in making informed clinical decisions.

However, the sparsity of EHR data typically allows for only a small fraction of medical codes to be learned effectively, thereby restricting the ability of deep learning approaches. To overcome this drawback, knowledge graphs have been applied to incorporate prior medical knowledge for these deep learning models, which augment the representation of medical codes to better support the downstream prediction tasks.

4.3.1 Intermediate Steps to Advance Prediction Models

ICD Coding aims to extract diagnosis and procedure codes from clinical notes, which often consisted of raw text [155–158]. It is often challenging, as the size of the candidate target codes can be large and the distribution of the codes is often long-tailed [159]. To overcome this, Xie et al. [160] and Cao et al. [161] propose to leverage knowledge graphs as *distant supervision* [54, 162], and inject the label information via structured *knowledge graph propagation* by leveraging graph convolution networks [163] to learn the correlations among medical codes. Besides, Lu et al. [164] propose to leverage knowledge graphs and the co-occurrence graph among clinical nodes simultaneously with a knowledge aggregation module to boost the ICD coding performance. Ren et al. [165] design a learning curriculum based on the hierarchical structure of the code to address the highly imbalanced label distribution issue and balance between frequent and rare labels. Overall, injecting additional knowledge with graph neural networks offers a way to mitigate the imbalanced label distribution issue and thus better.

Entity and Relation Extraction from Health Records. Health records contain rich unstructured or semi-structured data, making it difficult for clinicians to access and analyze relevant information. Entity and relation extraction helps convert this unstructured text into structured data that can be more easily processed, understood, and utilized. Specifically, *entity extraction* aims to identify entity mentions from clinical-free texts. There are two key steps for entity extraction, i.e., named entity recognition (NER) and disambiguation (NED). By leveraging additional knowledge graphs, Varma et al. [166] transfer structural knowledge from the knowledge base to the medical domain, improving rare entities’ disambiguation accuracy. Yuan et al. [167] inject additional knowledge from the knowledge graphs for entity linking and proposed two additional strategies, namely Post-pruning and Thresholding, to improve the efficiency and remove the effect of unlinkable entity mention. Fries et al. [69] leverage clinical ontologies to provide *weak supervision* sources to create additional training data for clinical entity disambiguation. Besides, *relation extraction* aims to identify and classify relationships between entities in unstructured text, which facilitates understanding complex biological processes, drug interactions, and disease mechanisms. To incorporate the external knowledge graph, several works [168, 169] proposed additional post-training steps to align the language models with biomedical

knowledge. Hong et al. [170] construct embeddings for a wide range of codified concepts from EHRs to identify relevant features related to a disease of interest, and Lin et al. [171] design a co-training scheme to jointly learn from text and knowledge graphs for extracting and classifying disease-disease relations. In summary, fusing knowledge graphs with language models can flexibly accommodate missing data types and bring additional performance gains, especially for those rare entities and relations.

4.3.2 Evidence Generation for Risk Prediction Models

Disease Prediction aims to predict the potential diseases of a given patient with his past clinical records. To assist the diagnosis with additional knowledge, GRAM [172] and KAME [173] utilize a medical ontology [174] where the leaf nodes are the medical codes found in EHR data, and their ancestors are more general categories. By incorporating information from medical ontologies into deep learning models via neural attention, these approaches learn better embeddings for different medical concepts to alleviate the data scarcity bottleneck. [175, 176] further consider the domain-specific knowledge graph KnowLife [177] to enrich the embeddings of medical entities with their neighbors on the knowledge graph. These approaches mainly directly update the embeddings of different concepts to improve the feature learning, but may be at the risk of ignoring the high-level order information from the knowledge graph. To tackle this drawback, Ye et al. [178] explicitly exploit *paths* in KG from the observed symptoms to the target disease to model the personalized information for diverse patients with a relational-guided attention mechanism. Xu et al. [179] design a self-supervised learning approach to pre-train a graph attention network for learning the embedding of medical concepts and completing the knowledge graph simultaneously. These approaches better harness the structure information, and often lead to better performance than the pure embedding-based knowledge integration techniques.

Treatment Recommendation aims to recommend personalized medications to patients based on their individual health conditions, which can help physicians select the most effective medications for their patients, and improve treatment outcomes [180–182]. To effectively exploit external knowledge, Shang et al. [183] use drug ontologies to design additional pretraining loss and directly improve the representation of drugs, and several studies [184, 185] attempt to extract the additional drug interaction graphs to model the negative side effects of specific drug pairs and reduce the possibility of recommending negative drug-drug interaction combinations. Besides, Wu et al. [186] leveraged ontologies to improve the drug representations, and facilitates drug recommendation under a more challenging few-shot setting.

4.4 Public Health

Public Health research can significantly benefit from HKGs. Knowledge graphs can help organize, structure, and formalize extensive information from diverse and heterogeneous sources. This allows researchers to analyze data, reason about factors, and make decisions on a larger scale.

Epidemiology. The field of epidemiology has seen an increase in the use of knowledge graphs to analyze and understand the spread of diseases. A study conducted by Gao et

al. [187] analyzes the research hotspots and development trends of wastewater-based epidemiology (WBE) using knowledge graphs constructed from nearly 900 papers. Domingo-Fernández et al. [188] create the COVID-19 Knowledge Graph, a comprehensive cause-and-effect network constructed from the scientific literature on the coronavirus. Additionally, Turki et al. [189] use knowledge graphs to assess and validate the portion of Wikidata related to COVID-19 epidemiology using an automatable task set. Pressat Laffouilhère et al. [190] develop OntoBioStat, a domain ontology related to covariate selection and bias in biostatistics, which can help interpret significant statistical associations between variables.

Environmental Health. Fecho et al. [191] develop ROBOKOP, a biomedical knowledge graph-based system, to validate associations between workplace chemical exposures and immune-mediated diseases. Wolffe et al. [192] propose using knowledge graphs in systematic evidence mapping in environmental health. This approach overcomes the limitations of rigid data tables by offering a more suitable model for handling the highly connected and complex nature of environmental health data.

Health Policy and Management. Wu et al. [193] have analyzed the COVID-19 epidemic situation using a knowledge graph of patient activity. This method enables in-depth study of the transmission process, analysis of key nodes, and tracing of activity tracks. Meanwhile, Yu et al. [194] develop a chronic management system, which combines knowledge graphs and big data to optimize the management of chronic diseases in children. This system enhances treatment and resource utilization while conforming to the requirements of the Chronic Care Model.

Social and Behavioral Health. There have been several interesting research studies. Cao et al. [195] build a high-level suicide-oriented knowledge graph combined with deep neural networks for detecting suicidal ideation on social media platforms. Also, Liu et al. [196] conduct a bibliometric analysis of driver behavior research. Additionally, Wang et al. [197] create an analysis framework for interpreting causal associations in emotional logic. They introduce a knowledge graph into appraisal theories, improving human emotional inference.

5 Promise and Outlook

The potential impact of comprehensive and fine-grained HKGs on biomedical research and clinical practice is significant. By integrating vast amounts of biomedical knowledge from multiple domains, HKGs can facilitate the discovery of new disease mechanisms and the identification of novel drug targets. They also help to enable personalized medicine by identifying patient subgroups with shared disease mechanisms. The recent success of large language models (*a.k.a.* foundation models) such as ChatGPT offers promising opportunities in capturing such semantics from the biomedical context [68, 198–200], enabling the construction of unprecedentedly comprehensive HKGs. In turn, HKGs also help improve LLMs by providing accurate and contextualized knowledge to regularize the generated content. This is particularly useful in evaluating LLMs in biomedical applications and addressing the problem of hallucination in critical areas.

6 Conclusion

Healthcare knowledge graphs have emerged as a promising approach for capturing and organizing medical knowledge in a structured and interpretable way. This comprehensive review paper provides an overview of the current state of HKGs, including their construction, utilization models, and applications in healthcare. Furthermore, the paper discusses the potential future developments of HKGs. In conclusion, HKGs have played a significant role in advancing health research. With the advent of LLMs, there are even more opportunities to combine HKGs and LLMs to reduce the generation of false or unreliable content. We hope that our comprehensive review of this field offers a helpful perspective for future reference.

References

- [1] Ji, S., Pan, S., Cambria, E., Marttinen, P., Philip, S.Y.: A survey on knowledge graphs: Representation, acquisition, and applications. *IEEE transactions on neural networks and learning systems* **33**(2), 494–514 (2021)
- [2] Nicholson, D.N., Greene, C.S.: Constructing knowledge graphs and their biomedical applications. *Computational and structural biotechnology journal* **18**, 1414–1428 (2020)
- [3] Wang, P., Jiang, H., Xu, J., Zhang, Q.: Knowledge graph construction and applications for web search and beyond. *Data Intelligence* **1**(4), 333–349 (2019)
- [4] Wang, X., He, X., Cao, Y., Liu, M., Chua, T.-S.: Kgat: Knowledge graph attention network for recommendation. In: *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 950–958 (2019)
- [5] Zhou, S., Dai, X., Chen, H., Zhang, W., Ren, K., Tang, R., He, X., Yu, Y.: Interactive recommender system via knowledge graph-enhanced reinforcement learning. In: *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 179–188 (2020)
- [6] Lin, B.Y., Chen, X., Chen, J., Ren, X.: Kagnet: Knowledge-aware graph networks for commonsense reasoning. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 2829–2839 (2019)
- [7] Yasunaga, M., Ren, H., Bosselut, A., Liang, P., Leskovec, J.: Qa-gnn: Reasoning with language models and knowledge graphs for question answering. In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 535–546 (2021)
- [8] Yan, J., Raman, M., Chan, A., Zhang, T., Rossi, R., Zhao, H., Kim, S., Lipka,

- N., Ren, X.: Learning contextualized knowledge structures for commonsense reasoning. In: Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021, pp. 4038–4051 (2021)
- [9] Kan, X., Cui, H., Yang, C.: Zero-shot scene graph relation prediction through commonsense knowledge integration. In: Machine Learning and Knowledge Discovery in Databases. Research Track: European Conference, ECML PKDD 2021, Bilbao, Spain, September 13–17, 2021, Proceedings, Part II 21, pp. 466–482 (2021). Springer
- [10] Santos, A., Colaço, A.R., Nielsen, A.B., Niu, L., Strauss, M., Geyer, P.E., Coscia, F., Albrechtsen, N.J.W., Mundt, F., Jensen, L.J., *et al.*: A knowledge graph to interpret clinical proteomics data. *Nature Biotechnology* **40**(5), 692–702 (2022)
- [11] Chandak, P., Huang, K., Zitnik, M.: Building a knowledge graph to enable precision medicine. *Nature Scientific Data* (2023) <https://doi.org/10.1038/s41597-023-01960-3>
- [12] Bouayad, L., Ialynytchev, A., Padmanabhan, B.: Patient health record systems scope and functionalities: literature review and future directions. *Journal of medical Internet research* **19**(11), 388 (2017)
- [13] Rajkomar, A., Oren, E., Chen, K., Dai, A.M., Hajaj, N., Hardt, M., Liu, P.J., Liu, X., Marcus, J., Sun, M., *et al.*: Scalable and accurate deep learning with electronic health records. *NPJ digital medicine* **1**(1), 18 (2018)
- [14] Mehta, N., Pandit, A.: Concurrence of big data analytics and healthcare: A systematic review. *International journal of medical informatics* **114**, 57–65 (2018)
- [15] Dash, S., Shakyawar, S.K., Sharma, M., Kaushik, S.: Big data in healthcare: management, analysis and future prospects. *Journal of Big Data* **6**(1), 1–25 (2019)
- [16] Devlin, J., Chang, M.-W., Lee, K., Toutanova, K.: Bert: Pre-training of deep bidirectional transformers for language understanding. In: Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers), pp. 4171–4186 (2019)
- [17] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J.D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., *et al.*: Language models are few-shot learners. *Advances in neural information processing systems* **33**, 1877–1901 (2020)
- [18] Zhang, J., Song, X., Zeng, Y., Chen, J., Shen, J., Mao, Y., Li, L.: Taxonomy completion via triplet matching network. In: Proceedings of the AAAI Conference

on Artificial Intelligence, vol. 35, pp. 4662–4670 (2021)

- [19] Yu, Y., Li, Y., Shen, J., Feng, H., Sun, J., Zhang, C.: Steam: Self-supervised taxonomy expansion with mini-paths. In: Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, pp. 1026–1035 (2020)
- [20] Wang, S., Zhao, R., Chen, X., Zheng, Y., Liu, B.: Enquire one’s parent and child before decision: Fully exploit hierarchical structure for self-supervised taxonomy expansion. In: Proceedings of the Web Conference 2021, pp. 3291–3304 (2021)
- [21] Zeng, Q., Lin, J., Yu, W., Cleland-Huang, J., Jiang, M.: Enhancing taxonomy completion with concept generation via fusing relational representations. In: Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining, pp. 2104–2113 (2021)
- [22] Xu, J., Kim, S., Song, M., Jeong, M., Kim, D., Kang, J., Rousseau, J.F., Li, X., Xu, W., Torvik, V.I., *et al.*: Building a pubmed knowledge graph. *Scientific data* **7**(1), 205 (2020)
- [23] Gyrard, A., Gaur, M., Shekarpour, S., Thirunarayan, K., Sheth, A.: Personalized health knowledge graph. In: CEUR Workshop Proceedings, vol. 2317 (2018)
- [24] Li, L., Wang, P., Yan, J., Wang, Y., Li, S., Jiang, J., Sun, Z., Tang, B., Chang, T.-H., Wang, S., *et al.*: Real-world data medical knowledge graph: construction and applications. *Artificial intelligence in medicine* **103**, 101817 (2020)
- [25] Hassani-Pak, K., Rawlings, C.: Knowledge discovery in biological databases for revealing candidate genes linked to complex phenotypes. *Journal of integrative bioinformatics* **14**(1) (2017)
- [26] Seneviratne, O., Harris, J., Chen, C.-H., McGuinness, D.L.: Personal health knowledge graph for clinically relevant diet recommendations. *arXiv preprint arXiv:2110.10131* (2021)
- [27] Eberhardt, J., Bilchik, A., Stojadinovic, A.: Clinical decision support systems: potential with pitfalls. *Journal of Surgical Oncology* **105**(5), 502–510 (2012)
- [28] Castaneda, C., Nalley, K., Mannion, C., Bhattacharyya, P., Blake, P., Pecora, A., Goy, A., Suh, K.S.: Clinical decision support systems for improving diagnostic accuracy and achieving precision medicine. *Journal of clinical bioinformatics* **5**(1), 1–16 (2015)
- [29] Guarino, N., Oberle, D., Staab, S.: What is an ontology? *Handbook on ontologies*, 1–17 (2009)
- [30] Blagec, K., Barbosa-Silva, A., Ott, S., Samwald, M.: A curated, ontology-based,

large-scale knowledge graph of artificial intelligence tasks and benchmarks. *Scientific Data* **9**(1), 322 (2022)

- [31] Ashburner, M., Ball, C.A., Blake, J.A., Botstein, D., Butler, H., Cherry, J.M., Davis, A.P., Dolinski, K., Dwight, S.S., Eppig, J.T., *et al.*: Gene ontology: tool for the unification of biology. *Nature genetics* **25**(1), 25–29 (2000)
- [32] Schriml, L.M., Arze, C., Nadendla, S., Chang, Y.-W.W., Mazaitis, M., Felix, V., Feng, G., Kibbe, W.A.: Disease ontology: a backbone for disease semantic integration. *Nucleic acids research* **40**(D1), 940–946 (2012)
- [33] Bard, J., Rhee, S.Y., Ashburner, M.: An ontology for cell types. *Genome biology* **6**(2), 1–5 (2005)
- [34] Song, B., Li, F., Liu, Y., Zeng, X.: Deep learning methods for biomedical named entity recognition: a survey and qualitative comparison. *Briefings in Bioinformatics* **22**(6), 282 (2021)
- [35] Xing, R., Luo, J., Song, T.: Biorel: towards large-scale biomedical relation extraction. *BMC bioinformatics* **21**, 1–13 (2020)
- [36] Hahn, U., Oleynik, M.: Medical information extraction in the age of deep learning. *Yearbook of medical informatics* **29**(01), 208–220 (2020)
- [37] Bodenreider, O.: The unified medical language system (umls): integrating biomedical terminology. *Nucleic acids research* **32**(suppl_1), 267–270 (2004)
- [38] Donnelly, K., *et al.*: Snomed-ct: The advanced terminology and coding system for ehealth. *Studies in health technology and informatics* **121**, 279 (2006)
- [39] Hirsch, J., Nicola, G., McGinty, G., Liu, R., Barr, R., Chittle, M., Manchikanti, L.: Icd-10: history and context. *American Journal of Neuroradiology* **37**(4), 596–599 (2016)
- [40] Wang, P., Shi, T., Reddy, C.K.: Text-to-sql generation for question answering on electronic medical records. In: *Proceedings of The Web Conference 2020*, pp. 350–361 (2020)
- [41] Bordes, A., Usunier, N., Garcia-Duran, A., Weston, J., Yakhnenko, O.: Translating embeddings for modeling multi-relational data. *Advances in neural information processing systems* **26** (2013)
- [42] Lu, J., Yang, C.: Open-world taxonomy and knowledge graph co-learning. In: *4th Conference on Automated Knowledge Base Construction. AKBC 2022* (2022)
- [43] Himmelstein, D.S., Lizée, A., Hessler, C., Brueggeman, L., Chen, S.L., Hadley, D., Green, A., Khankhanian, P., Baranzini, S.E.: Systematic integration of biomedical knowledge prioritizes drugs for repurposing. *Elife* **6**, 26726 (2017)

- [44] Su, C., Hou, Y., Zhou, M., Rajendran, S., Maasch, J.R., Abedi, Z., Zhang, H., Bai, Z., Cuturrufo, A., Guo, W., et al.: Biomedical discovery through the integrative biomedical knowledge hub (ibkh). *Iscience* (2023)
- [45] Youn, J., Rai, N., Tagkopoulos, I.: Knowledge integration and decision support for accelerated discovery of antibiotic resistance genes. *Nature Communications* **13**(1), 2360 (2022)
- [46] Faria, D., Jiménez-Ruiz, E., Pesquita, C., Santos, E., Couto, F.M.: Towards annotating potential incoherences in bioportal mappings. In: *The Semantic Web–ISWC 2014: 13th International Semantic Web Conference*, Riva del Garda, Italy, October 19–23, 2014. *Proceedings, Part II* 13, pp. 17–32 (2014). Springer
- [47] He, Y., Chen, J., Dong, H., Jiménez-Ruiz, E., Hadian, A., Horrocks, I.: Machine learning-friendly biomedical datasets for equivalence and subsumption ontology matching. In: *The Semantic Web–ISWC 2022: 21st International Semantic Web Conference*, Virtual Event, October 23–27, 2022, *Proceedings*, pp. 575–591 (2022). Springer
- [48] Suchanek, F.M., Abiteboul, S., Senellart, P.: Paris: probabilistic alignment of relations, instances, and schema. *Proceedings of the VLDB Endowment* **5**(3), 157–168 (2011)
- [49] Maaroufi, M., Choquet, R., Landais, P., Jaulent, M.-C.: Formalizing mappings to optimize automated schema alignment: application to rare diseases. In: *e-Health–For Continuity of Care*, pp. 283–287. IOS Press, ??? (2014)
- [50] Bachman, J.A., Gyori, B.M., Sorger, P.K.: Famplex: a resource for entity recognition and relationship resolution of human protein families and complexes in biomedical text mining. *BMC bioinformatics* **19**, 1–14 (2018)
- [51] Hu, Y., Nie, T., Shen, D., Kou, Y., Yu, G.: An integrated pipeline model for biomedical entity alignment. *Frontiers of Computer Science* **15**, 1–15 (2021)
- [52] Ma, J., Zhou, C., Chen, Y., Wang, Y., Hu, G., Qiao, Y.: Tetre: A novel temporal conflict resolution method based on temporal knowledge graph embedding. *Information* **14**(3), 155 (2023)
- [53] Ye, H., Zhang, N., Chen, H., Chen, H.: Generative knowledge graph construction: A review. In: *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 1–17 (2022)
- [54] Liang, C., Yu, Y., Jiang, H., Er, S., Wang, R., Zhao, T., Zhang, C.: Bond: Bert-assisted open-domain named entity recognition with distant supervision. In: *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 1054–1064 (2020)

- [55] Chen, X., Li, L., Fei, Q., Zhang, N., Tan, C., Jiang, Y., Huang, F., Chen, H.: One model for all domains: Collaborative domain-prefix tuning for cross-domain ner. arXiv preprint arXiv:2301.10410 (2023)
- [56] Huang, Y., He, K., Wang, Y., Zhang, X., Gong, T., Mao, R., Li, C.: Copner: Contrastive learning with prompt guiding for few-shot named entity recognition. In: Proceedings of the 29th International Conference on Computational Linguistics, pp. 2515–2527 (2022)
- [57] Liu, A.T., Xiao, W., Zhu, H., Zhang, D., Li, S.-W., Arnold, A.: Qaner: Prompting question answering models for few-shot named entity recognition. arXiv preprint arXiv:2203.01543 (2022)
- [58] Zhuang, Y., Li, Y., Zhang, J., Yu, Y., Mou, Y., Chen, X., Song, L., Zhang, C.: ReSel: N-ary relation extraction from scientific text and tables by learning to retrieve and select. In: Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, pp. 730–744 (2022)
- [59] Lu, Y., Liu, Q., Dai, D., Xiao, X., Lin, H., Han, X., Sun, L., Wu, H.: Unified structure generation for universal information extraction. In: Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp. 5755–5772 (2022)
- [60] Yang, X., Chen, A., PourNejatian, N., Shin, H.C., Smith, K.E., Parisien, C., Compas, C., Martin, C., Costa, A.B., Flores, M.G., *et al.*: A large language model for electronic health records. npj Digital Medicine **5**(1), 194 (2022)
- [61] De Cao, N., Wu, L., Popat, K., Artetxe, M., Goyal, N., Plekhanov, M., Zettlemoyer, L., Cancedda, N., Riedel, S., Petroni, F.: Multilingual autoregressive entity linking. Transactions of the Association for Computational Linguistics **10**, 274–290
- [62] Mrini, K., Nie, S., Gu, J., Wang, S., Sanjabi, M., Firooz, H.: Detection, disambiguation, re-ranking: Autoregressive entity linking as a multi-task problem. In: Findings of the Association for Computational Linguistics: ACL 2022, pp. 1972–1983 (2022)
- [63] Cho, Y.M., Zhang, L., Callison-Burch, C.: Unsupervised entity linking with guided summarization and multiple-choice selection. In: Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, pp. 9394–9401 (2022)
- [64] Saxena, A., Kochsiek, A., Gemulla, R.: Sequence-to-sequence knowledge graph completion and question answering. In: Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp. 2814–2828 (2022)

- [65] Xie, X., Zhang, N., Li, Z., Deng, S., Chen, H., Xiong, F., Chen, M., Chen, H.: From discrimination to generation: knowledge graph completion with generative transformer. In: Companion Proceedings of the Web Conference 2022, pp. 162–165 (2022)
- [66] Shen, J., Wang, C., Yuan, Y., Han, J., Ji, H., Sen, K., Zhang, M., Song, D.: PALT: Parameter-lite transfer of language models for knowledge graph completion. In: Findings of the Association for Computational Linguistics: EMNLP 2022, pp. 3833–3847 (2022)
- [67] Zhang, J., Wang, Z., Zhang, S., Bhalerao, M.M., Liu, Y., Zhu, D., Wang, S.: Graphprompt: Biomedical entity normalization using graph-based prompt templates. In: The 37th AAAI Conference on Artificial Intelligence (2023)
- [68] Agrawal, M., Hegselmann, S., Lang, H., Kim, Y., Sontag, D.: Large language models are few-shot clinical information extractors. In: Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, pp. 1998–2022 (2022)
- [69] Fries, J.A., Steinberg, E., Khattar, S., Fleming, S.L., Posada, J., Callahan, A., Shah, N.H.: Ontology-driven weak supervision for clinical entity classification in electronic health records. *Nature communications* **12**(1), 2017 (2021)
- [70] Hu, Y., Ameer, I., Zuo, X., Peng, X., Zhou, Y., Li, Z., Li, Y., Li, J., Jiang, X., Xu, H.: Zero-shot clinical entity recognition using chatgpt. arXiv preprint arXiv:2303.16416 (2023)
- [71] Zhu, T., Qin, Y., Chen, Q., Hu, B., Xiang, Y.: Enhancing entity representations with prompt learning for biomedical entity linking. In: Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence, IJCAI 2022, pp. 4036–4042 (2022)
- [72] Lu, J., Shen, J., Xiong, B., Ma, W., Steffen, S., Yang, C.: Hiprompt: Few-shot biomedical knowledge fusion via hierarchy-oriented prompting. In: Proceedings of The 46th International ACM SIGIR Conference on Research and Development in Information Retrieval - Short Paper. SIGIR 2023 (2023)
- [73] Shi, B., Weninger, T.: Open-world knowledge graph completion. In: Proceedings of the AAAI Conference on Artificial Intelligence, vol. 32 (2018)
- [74] Das, R., Godbole, A., Monath, N., Zaheer, M., McCallum, A.: Probabilistic case-based reasoning for open-world knowledge graph completion. In: Findings of the Association for Computational Linguistics: EMNLP 2020, pp. 4752–4765 (2020)
- [75] Niu, L., Fu, C., Yang, Q., Li, Z., Chen, Z., Liu, Q., Zheng, K.: Open-world knowledge graph completion with multiple interaction attention. *World Wide Web* **24**, 419–439 (2021)

- [76] Li, Q., Wang, D., Song, S.F.K., Zhang, Y., Yu, G.: Oerl: Enhanced representation learning via open knowledge graphs. *IEEE Transactions on Knowledge and Data Engineering*, 1–14 (2022) <https://doi.org/10.1109/TKDE.2022.3218850>
- [77] Himmelstein, D.S., Baranzini, S.E.: Heterogeneous network edge prediction: a data integration approach to prioritize disease-associated genes. *PLoS computational biology* **11**(7), 1004259 (2015)
- [78] Ioannidis, V.N., Song, X., Manchanda, S., Li, M., Pan, X., Zheng, D., Ning, X., Zeng, X., Karypis, G.: DRKG - Drug Repurposing Knowledge Graph for Covid-19. <https://github.com/gnn4dr/DRKG/> (2020)
- [79] Ranjan, A., Shukla, S., Datta, D., Misra, R.: Generating novel molecule for target protein (sars-cov-2) using drug–target interaction based on graph neural network. *Network Modeling Analysis in Health Informatics and Bioinformatics* **11**, 1–11 (2022)
- [80] Li, B., Chen, H.: Prediction of compound synthesis accessibility based on reaction knowledge graph. *Molecules* **27**(3), 1039 (2022)
- [81] Jeong, J., Lee, N., Shin, Y., Shin, D.: Intelligent generation of optimal synthetic pathways based on knowledge graph inference and retrosynthetic predictions using reaction big data. *Journal of the Taiwan Institute of Chemical Engineers* **130**, 103982 (2022)
- [82] Zhang, R., Hristovski, D., Schutte, D., Kastrin, A., Fiszman, M., Kilicoglu, H.: Drug repurposing for covid-19 via knowledge graph completion. *Journal of biomedical informatics* **115**, 103696 (2021)
- [83] Gao, Z., Ding, P., Xu, R.: Kg-predict: a knowledge graph computational framework for drug repurposing. *Journal of biomedical informatics* **132**, 104133 (2022)
- [84] Zhang, X., Che, C.: Drug repurposing for parkinson’s disease by integrating knowledge graph completion model and knowledge fusion of medical literature. *Future Internet* **13**(1), 14 (2021)
- [85] Ghorbanali, Z., Zare-Mirakabad, F., Akbari, M., Salehi, N., Masoudi-Nejad, A.: Drugrep-kg: Toward learning a unified latent space for drug repurposing using knowledge graphs. *Journal of Chemical Information and Modeling* (2023)
- [86] Wishart, D.S., Feunang, Y.D., Guo, A.C., Lo, E.J., Marcu, A., Grant, J.R., Sajed, T., Johnson, D., Li, C., Sayeeda, Z., *et al.*: Drugbank 5.0: a major update to the drugbank database for 2018. *Nucleic acids research* **46**(D1), 1074–1082 (2018)
- [87] Ernst, P., Meng, C., Siu, A., Weikum, G.: Knowlife: a knowledge graph for

- health and life sciences. In: 2014 IEEE 30th International Conference on Data Engineering, pp. 1254–1257 (2014). IEEE
- [88] Zheng, S., Rao, J., Song, Y., Zhang, J., Xiao, X., Fang, E.F., Yang, Y., Niu, Z.: Pharmkg: a dedicated knowledge graph benchmark for biomedical data mining. *Briefings in bioinformatics* **22**(4), 344 (2021)
 - [89] Bizon, C., Cox, S., Balhoff, J., Kebede, Y., Wang, P., Morton, K., Fecho, K., Tropsha, A.: Robokop kg and kgb: integrated knowledge graphs from federated sources. *Journal of chemical information and modeling* **59**(12), 4968–4973 (2019)
 - [90] Kim, D., Bae, S., Kim, S., Choi, E.: Uncertainty-aware text-to-program for question answering on structured electronic health records. In: Flores, G., Chen, G.H., Pollard, T., Ho, J.C., Naumann, T. (eds.) *Proceedings of the Conference on Health, Inference, and Learning*. *Proceedings of Machine Learning Research*, vol. 174, pp. 138–151. PMLR, ??? (2022)
 - [91] Liu, Q., Yogatama, D., Blunsom, P.: Relational memory-augmented language models. *Transactions of the Association for Computational Linguistics* **10**, 555–572 (2022)
 - [92] Guu, K., Lee, K., Tung, Z., Pasupat, P., Chang, M.: Retrieval augmented language model pre-training. In: *International Conference on Machine Learning*, pp. 3929–3938 (2020). PMLR
 - [93] Xu, R., Yu, Y., Ho, J.C., Yang, C.: Weakly-supervised scientific document classification via retrieval-augmented multi-stage training. In: *the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval* (2023)
 - [94] Shi, W., Min, S., Yasunaga, M., Seo, M., James, R., Lewis, M., Zettlemoyer, L., Yih, W.-t.: Replug: Retrieval-augmented black-box language models. *arXiv preprint arXiv:2301.12652* (2023)
 - [95] Tchechmedjiev, A., Fafalios, P., Boland, K., Gasquet, M., Zloch, M., Zapilko, B., Dietze, S., Todorov, K.: Claimskg: A knowledge graph of fact-checked claims. In: *The Semantic Web–ISWC 2019: 18th International Semantic Web Conference, Auckland, New Zealand, October 26–30, 2019, Proceedings, Part II* 18, pp. 309–324 (2019). Springer
 - [96] Vedula, N., Parthasarathy, S.: Face-keg: Fact checking explained using knowledge graphs. In: *Proceedings of the 14th ACM International Conference on Web Search and Data Mining*, pp. 526–534 (2021)
 - [97] Mayank, M., Sharma, S., Sharma, R.: Deap-faked: Knowledge graph based approach for fake news detection. In: *2022 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*, pp. 47–51

(2022). IEEE

- [98] Shirai, S.S., Seneviratne, O., Gordon, M.E., Chen, C.-H., McGuinness, D.L.: Identifying ingredient substitutions using a knowledge graph of food. *Frontiers in Artificial Intelligence* **3**, 621766 (2021)
- [99] Mengoni, P., Yang, J.: Empowering covid-19 fact-checking with extended knowledge graphs. In: *Computational Science and Its Applications–ICCSA 2022 Workshops: Malaga, Spain, July 4–7, 2022, Proceedings, Part I*, pp. 138–150 (2022). Springer
- [100] Yu, Y., Huang, K., Zhang, C., Glass, L.M., Sun, J., Xiao, C.: Sumgnn: multi-typed drug interaction prediction via efficient knowledge graph summarization. *Bioinformatics* **37**(18), 2988–2995 (2021)
- [101] Su, X., Hu, L., You, Z., Hu, P., Zhao, B.: Attention-based knowledge graph representation learning for predicting drug-drug interactions. *Briefings in bioinformatics* **23**(3), 140 (2022)
- [102] Muggleton, S.: *Inductive Logic Programming* vol. 38. Morgan Kaufmann, ??? (1992)
- [103] Galárraga, L.A., Teflioudi, C., Hose, K., Suchanek, F.: Amie: association rule mining under incomplete evidence in ontological knowledge bases. In: *Proceedings of the 22nd International Conference on World Wide Web*, pp. 413–422 (2013)
- [104] Kok, S., Domingos, P.: Learning the structure of markov logic networks. In: *Proceedings of the 22nd International Conference on Machine Learning*, pp. 441–448 (2005)
- [105] Alshahrani, M., Khan, M.A., Maddouri, O., Kinjo, A.R., Queralt-Rosinach, N., Hoehndorf, R.: Neuro-symbolic representation learning on biological knowledge graphs. *Bioinformatics* **33**(17), 2723–2730 (2017)
- [106] Zhu, Z., Galkin, M., Zhang, Z., Tang, J.: Neural-symbolic models for logical queries on knowledge graphs. In: *International Conference on Machine Learning*, pp. 27454–27478 (2022). PMLR
- [107] Giacomini, K.M., Krauss, R.M., Roden, D.M., Eichelbaum, M., Hayden, M.R., Nakamura, Y.: When good drugs go bad. *Nature* **446**(7139), 975–977 (2007)
- [108] Karim, M.R., Cochez, M., Jares, J.B., Uddin, M., Beyan, O., Decker, S.: Drug-drug interaction prediction based on knowledge graph embeddings and convolutional-lstm network. In: *Proceedings of the 10th ACM International Conference on Bioinformatics, Computational Biology and Health Informatics*, pp. 113–123 (2019)

- [109] Dai, Y., Guo, C., Guo, W., Eickhoff, C.: Drug–drug interaction prediction with wasserstein adversarial autoencoder-based knowledge graph embeddings. *Briefings in bioinformatics* **22**(4), 256 (2021)
- [110] Lin, X., Quan, Z., Wang, Z.-J., Ma, T., Zeng, X.: Kggn: Knowledge graph neural network for drug-drug interaction prediction. In: *IJCAI*, vol. 380, pp. 2739–2745 (2020)
- [111] Chen, X., Yan, C.C., Zhang, X., Zhang, X., Dai, F., Yin, J., Zhang, Y.: Drug–target interaction prediction: databases, web servers and computational models. *Briefings in bioinformatics* **17**(4), 696–712 (2016)
- [112] Li, G., Sun, W., Xu, J., Hu, L., Zhang, W., Zhang, P.: Ga-ens: A novel drug–target interactions prediction method by incorporating prior knowledge graph into dual wasserstein generative adversarial network with gradient penalty. *Applied Soft Computing* **139**, 110151 (2023)
- [113] Zhang, S., Lin, X., Zhang, X.: Discovering dti and ddi by knowledge graph with mhrw and improved neural network. In: *2021 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pp. 588–593 (2021). IEEE
- [114] Wang, S., Du, Z., Ding, M., Rodriguez-Paton, A., Song, T.: Kg-dti: a knowledge graph based deep learning method for drug-target interaction predictions and alzheimer’s disease drug repositions. *Applied Intelligence* **52**(1), 846–857 (2022)
- [115] Ye, Q., Hsieh, C.-Y., Yang, Z., Kang, Y., Chen, J., Cao, D., He, S., Hou, T.: A unified drug–target interaction prediction framework based on knowledge graph and recommendation system. *Nature communications* **12**(1), 6775 (2021)
- [116] Zhang, B., Fu, Y., Lu, Y., Zhang, Z., Clarke, R., Van Eyk, J.E., Herrington, D.M., Wang, Y.: DDN2.0: R and python packages for differential dependency network analysis of biological systems. *bioRxiv*, 2021–04 (2021)
- [117] Quan, X., Cai, W., Xi, C., Wang, C., Yan, L.: AIMedGraph: A comprehensive multi-relational knowledge graph for precision medicine **2023**, 006
- [118] Linnarsson, S., Teichmann, S.A.: *Single-cell genomics: coming of age*. Springer (2016)
- [119] Fang, L., Li, Y., Ma, L., Xu, Q., Tan, F., Chen, G.: Grndb: decoding the gene regulatory networks in diverse human and mouse conditions. *Nucleic acids research* **49**(D1), 97–103 (2021)
- [120] Feng, F., Tang, F., Gao, Y., Zhu, D., Li, T., Yang, S., Yao, Y., Huang, Y., Liu, J.: Genomickb: a knowledge graph for the human genome. *Nucleic Acids Research* **51**(D1), 950–956 (2023)

- [121] Eltzschig, H.K., Weissmüller, T., Mager, A., Eckle, T.: Nucleotide metabolism and cell-cell interactions. *Cell-Cell Interactions: Methods and Protocols*, 73–87 (2006)
- [122] Efremova, M., Vento-Tormo, M., Teichmann, S.A., Vento-Tormo, R.: Cellphonedb: inferring cell–cell communication from combined expression of multi-subunit ligand–receptor complexes. *Nature protocols* **15**(4), 1484–1506 (2020)
- [123] Jin, S., Guerrero-Juarez, C.F., Zhang, L., Chang, I., Ramos, R., Kuan, C.-H., Myung, P., Plikus, M.V., Nie, Q.: Inference and analysis of cell-cell communication using cellchat. *Nature communications* **12**(1), 1088 (2021)
- [124] Liu, Z., Sun, D., Wang, C.: Evaluation of cell-cell interaction methods by integrating single-cell rna sequencing data with spatial information. *Genome Biology* **23**(1), 1–38 (2022)
- [125] Dimitrov, D., Türei, D., Garrido-Rodriguez, M., Burmedi, P.L., Nagai, J.S., Boys, C., Ramirez Flores, R.O., Kim, H., Szalai, B., Costa, I.G., *et al.*: Comparison of methods and resources for cell-cell communication inference from single-cell rna-seq data. *Nature Communications* **13**(1), 3224 (2022)
- [126] Shao, X., Liao, J., Li, C., Lu, X., Cheng, J., Fan, X.: Celltalkdb: a manually curated database of ligand–receptor interactions in humans and mice. *Briefings in bioinformatics* **22**(4), 269 (2021)
- [127] Shao, X., Li, C., Yang, H., Lu, X., Liao, J., Qian, J., Wang, K., Cheng, J., Yang, P., Chen, H., *et al.*: Knowledge-graph-based cell-cell communication inference for spatially resolved transcriptomic data with spatalk. *Nature Communications* **13**(1), 4429 (2022)
- [128] Adams, C.P., Brantner, V.V.: Estimating the cost of new drug development: is it really \$802 million? *Health affairs* **25**(2), 420–428 (2006)
- [129] DiMasi, J.A., Grabowski, H.G., Hansen, R.W.: The cost of drug development. *New England Journal of Medicine* **372**(20), 1972–1972 (2015)
- [130] Du, Y., Wang, S., Guo, X., Cao, H., Hu, S., Jiang, J., Varala, A., Angirekula, A., Zhao, L.: Graphgt: Machine learning datasets for graph generation and transformation. In: *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)* (2021)
- [131] Zeng, X., Tu, X., Liu, Y., Fu, X., Su, Y.: Toward better drug discovery with knowledge graph. *Current opinion in structural biology* **72**, 114–126 (2022)
- [132] Wang, S., Guo, X., Lin, X., Pan, B., Du, Y., Wang, Y., Ye, Y., Petersen, A., Leitgeb, A., AlKhalifa, S., *et al.*: Multi-objective deep data generation with correlated property control. *Advances in Neural Information Processing Systems*

- [133] Pan, B., Wang, Y., Lin, X., Qin, M., Du, Y., Ghaemi, S., Ding, A., Wang, S., Alkhalifa, S., Minbiole, K., *et al.*: Property-controllable generation of quaternary ammonium compounds. In: 2022 IEEE International Conference on Bioinformatics and Biomedicine (BIBM), pp. 3462–3469 (2022). IEEE
- [134] Wang, S., Du, Y., Guo, X., Pan, B., Qin, Z., Zhao, L.: Controllable data generation by deep learning: A review. arXiv preprint arXiv:2207.09542 (2022)
- [135] Nicolaou, C.A., Brown, N., Pattichis, C.S.: Molecular optimization using computational multi-objective methods. *Current Opinion in Drug Discovery and Development* **10**(3), 316 (2007)
- [136] Nicolaou, C.A., Pattichis, C.S.: Molecular substructure mining approaches for computer-aided drug discovery: A review. *Proc. of ITAB*, 26–28 (2006)
- [137] Jing, Y., Bian, Y., Hu, Z., Wang, L., Xie, X.-Q.S.: Deep learning for drug design: an artificial intelligence paradigm for drug discovery in the big data era. *The AAPS journal* **20**, 1–10 (2018)
- [138] Fu, T., Gao, W., Xiao, C., Yasonik, J., Coley, C.W., Sun, J.: Differentiable scaffolding tree for molecular optimization. arXiv preprint arXiv:2109.10469 (2021)
- [139] Fu, T., Sun, J.: Antibody complementarity determining regions (cdrs) design using constrained energy model. In: *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 389–399 (2022)
- [140] Pan, X., Lin, X., Cao, D., Zeng, X., Yu, P.S., He, L., Nussinov, R., Cheng, F.: Deep learning for drug repurposing: Methods, databases, and applications. *Wiley interdisciplinary reviews: Computational molecular science* **12**(4), 1597 (2022)
- [141] Huang, K., Fu, T., Glass, L.M., Zitnik, M., Xiao, C., Sun, J.: Deeppurpose: a deep learning library for drug–target interaction prediction. *Bioinformatics* **36**(22-23), 5545–5547 (2020)
- [142] Hasöksüz, M., Kilic, S., Sarac, F.: Coronaviruses and sars-cov-2. *Turkish journal of medical sciences* **50**(9), 549–556 (2020)
- [143] Zhu, Y., Che, C., Jin, B., Zhang, N., Su, C., Wang, F.: Knowledge-driven drug repurposing using a comprehensive drug knowledge graph. *Health Informatics Journal* **26**(4), 2737–2750 (2020)
- [144] MacLean, F.: Knowledge graphs and their applications in drug discovery. *Expert opinion on drug discovery* **16**(9), 1057–1069 (2021)
- [145] Xu, B., Liu, Y., Yu, S., Wang, L., Dong, J., Lin, H., Yang, Z., Wang, J., Xia,

- F.: A network embedding model for pathogenic genes prediction by multi-path random walking on heterogeneous network. *BMC Medical Genomics* **12**, 1–12 (2019)
- [146] Martin, L., Hutchens, M., Hawkins, C., Radnov, A.: How much do clinical trials cost. *Nat Rev Drug Discov* **16**(6), 381–382 (2017)
- [147] Rivera, S.C., Liu, X., Chan, A.-W., Denniston, A.K., Calvert, M.J., Ashrafian, H., Beam, A.L., Collins, G.S., Darzi, A., Deeks, J.J., *et al.*: Guidelines for clinical trial protocols for interventions involving artificial intelligence: the spirit-ai extension. *The Lancet Digital Health* **2**(10), 549–560 (2020)
- [148] He, Z., Tang, X., Yang, X., Guo, Y., George, T.J., Charness, N., Quan Hem, K.B., Hogan, W., Bian, J.: Clinical trial generalizability assessment in the big data era: a review. *Clinical and translational science* **13**(4), 675–684 (2020)
- [149] Yuan, C., Ryan, P.B., Ta, C., Guo, Y., Li, Z., Hardin, J., Makadia, R., Jin, P., Shang, N., Kang, T., *et al.*: Criteria2query: a natural language interface to clinical databases for cohort definition. *Journal of the American Medical Informatics Association* **26**(4), 294–305 (2019)
- [150] Tseo, Y., Salkola, M., Mohamed, A., Kumar, A., Abnousi, F.: Information extraction of clinical trial eligibility criteria. *arXiv preprint arXiv:2006.07296* (2020)
- [151] Liu, R., Rizzo, S., Whipple, S., Pal, N., Pineda, A.L., Lu, M., Arnieri, B., Lu, Y., Capra, W., Copping, R., *et al.*: Evaluating eligibility criteria of oncology trials using real-world data and ai. *Nature* **592**(7855), 629–633 (2021)
- [152] Gao, J., Xiao, C., Glass, L.M., Sun, J.: Compose: cross-modal pseudo-siamese network for patient trial matching. In: *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 803–812 (2020)
- [153] Fu, T., Huang, K., Xiao, C., Glass, L.M., Sun, J.: Hint: Hierarchical interaction network for clinical-trial-outcome predictions. *Patterns* **3**(4), 100445 (2022)
- [154] Wang, Z., Xiao, C., Sun, J.: Spot: Sequential predictive modeling of clinical trial outcome with meta-learning. *arXiv preprint arXiv:2304.05352* (2023)
- [155] Mullenbach, J., Wiegrefe, S., Duke, J., Sun, J., Eisenstein, J.: Explainable prediction of medical codes from clinical text. In: *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 1101–1111 (2018)
- [156] Zhang, Z., Liu, J., Razavian, N.: Bert-xml: Large scale automated icd coding

- using bert pretraining. In: Proceedings of the 3rd Clinical Natural Language Processing Workshop, pp. 24–34 (2020)
- [157] Vu, T., Nguyen, D.Q., Nguyen, A.: A label attention model for icd coding from clinical text. In: Proceedings of the Twenty-Ninth International Joint Conferences on Artificial Intelligence, pp. 3335–3341 (2021)
 - [158] Dong, H., Falis, M., Whiteley, W., Alex, B., Matterson, J., Ji, S., Chen, J., Wu, H.: Automated clinical coding: what, why, and where we are? NPJ digital medicine **5**(1), 159 (2022)
 - [159] Kim, B.-H., Ganapathi, V.: Read, attend, and code: Pushing the limits of medical codes prediction from clinical notes by machines. In: Proceedings of the 6th Machine Learning for Healthcare Conference. Proceedings of Machine Learning Research, vol. 149, pp. 196–208 (2021)
 - [160] Xie, X., Xiong, Y., Yu, P.S., Zhu, Y.: Ehr coding with multi-scale feature attention and structured knowledge graph propagation. In: Proceedings of the 28th ACM International Conference on Information and Knowledge Management, pp. 649–658 (2019)
 - [161] Cao, P., Chen, Y., Liu, K., Zhao, J., Liu, S., Chong, W.: Hypercore: Hyperbolic and co-graph representation for automatic icd coding. In: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pp. 3105–3114 (2020)
 - [162] Min, B., Grishman, R., Wan, L., Wang, C., Gondek, D.: Distant supervision for relation extraction with an incomplete knowledge base. In: Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, pp. 777–782 (2013)
 - [163] Kipf, T.N., Welling, M.: Semi-supervised classification with graph convolutional networks. The International Conference on Learning Representations (ICLR) (2016)
 - [164] Lu, J., Du, L., Liu, M., Dipnall, J.: Multi-label few/zero-shot learning with knowledge aggregated from multiple label graphs. In: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 2935–2943 (2020)
 - [165] Ren, W., Zeng, R., Wu, T., Zhu, T., Krishnan, R.G.: Hicu: Leveraging hierarchy for curriculum learning in automated icd coding. In: Proceedings of the 7th Machine Learning for Healthcare Conference. Proceedings of Machine Learning Research, vol. 182, pp. 198–223 (2022)
 - [166] Varma, M., Orr, L., Wu, S., Leszczynski, M., Ling, X., Ré, C.: Cross-domain data integration for named entity disambiguation in biomedical text. In: Findings

- of the Association for Computational Linguistics: EMNLP 2021, pp. 4566–4575 (2021)
- [167] Yuan, H., Lu, K., Yuan, Z.: Exploring partial knowledge base inference in biomedical entity linking. arXiv preprint arXiv:2303.10330 (2023)
 - [168] Fei, H., Ren, Y., Zhang, Y., Ji, D., Liang, X.: Enriching contextualized language model from knowledge graph for biomedical information extraction. *Briefings in bioinformatics* **22**(3), 110 (2021)
 - [169] Roy, A., Pan, S.: Incorporating medical knowledge in bert for clinical relation extraction. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 5357–5366 (2021)
 - [170] Hong, C., Rush, E., Liu, M., Zhou, D., Sun, J., Sonabend, A., Castro, V.M., Schubert, P., Panickan, V.A., Cai, T., *et al.*: Clinical knowledge extraction via sparse embedding regression (keser) with multi-center large scale electronic health record data. *NPJ digital medicine* **4**(1), 151 (2021)
 - [171] Lin, Y., Lu, K., Yu, S., Cai, T., Zitnik, M.: Multimodal learning on graphs for disease relation extraction. arXiv preprint arXiv:2203.08893 (2022)
 - [172] Choi, E., Bahadori, M.T., Song, L., Stewart, W.F., Sun, J.: Gram: graph-based attention model for healthcare representation learning. In: *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 787–795 (2017)
 - [173] Ma, F., You, Q., Xiao, H., Chitta, R., Zhou, J., Gao, J.: Kame: Knowledge-based attention model for diagnosis prediction in healthcare. In: *Proceedings of the 27th ACM International Conference on Information and Knowledge Management*, pp. 743–752 (2018)
 - [174] Dubberke, E.R., Reske, K.A., McDonald, L.C., Fraser, V.J.: Icd-9 codes and surveillance for clostridium difficile-associated disease. *Emerging infectious diseases* **12**(10), 1576 (2006)
 - [175] Yin, C., Zhao, R., Qian, B., Lv, X., Zhang, P.: Domain knowledge guided deep learning with electronic health records. In: *2019 IEEE International Conference on Data Mining (ICDM)*, pp. 738–747 (2019). IEEE
 - [176] Zhang, X., Qian, B., Li, Y., Yin, C., Wang, X., Zheng, Q.: Knowrisk: an interpretable knowledge-guided model for disease risk prediction. In: *2019 IEEE International Conference on Data Mining (ICDM)*, pp. 1492–1497 (2019). IEEE
 - [177] Ernst, P., Siu, A., Weikum, G.: Knowlife: a versatile approach for constructing a large knowledge graph for biomedical sciences. *BMC bioinformatics* **16**, 1–13 (2015)

- [178] Ye, M., Cui, S., Wang, Y., Luo, J., Xiao, C., Ma, F.: Medpath: Augmenting health risk prediction via medical knowledge paths. In: Proceedings of the Web Conference 2021, pp. 1397–1409 (2021)
- [179] Xu, X., Xu, X., Sun, Y., Liu, X., Li, X., Xie, G., Wang, F.: Predictive modeling of clinical events with mutual enhancement between longitudinal patient records and medical knowledge graph. In: 2021 IEEE International Conference on Data Mining (ICDM), pp. 777–786 (2021). IEEE
- [180] Zhang, Y., Chen, R., Tang, J., Stewart, W.F., Sun, J.: Leap: learning to prescribe effective and safe treatment combinations for multimorbidity. In: Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 1315–1324 (2017)
- [181] Bhoi, S., Lee, M.L., Hsu, W., Fang, H.S.A., Tan, N.C.: Personalizing medication recommendation with a graph-based approach. *ACM Transactions on Information Systems (TOIS)* **40**(3), 1–23 (2021)
- [182] Shang, J., Xiao, C., Ma, T., Li, H., Sun, J.: Gamenet: Graph augmented memory networks for recommending medication combination. In: Proceedings of the AAAI Conference on Artificial Intelligence, vol. 33, pp. 1126–1133 (2019)
- [183] Shang, J., Ma, T., Xiao, C., Sun, J.: Pre-training of graph augmented transformers for medication recommendation. In: Proceedings of the 28th International Joint Conference on Artificial Intelligence, pp. 5953–5959 (2019)
- [184] Wu, J., Qian, B., Li, Y., Gao, Z., Ju, M., Yang, Y., Zheng, Y., Gong, T., Li, C., Zhang, X.: Leveraging multiple types of domain knowledge for safe and effective drug recommendation. In: Proceedings of the 31st ACM International Conference on Information & Knowledge Management, pp. 2169–2178 (2022)
- [185] Tan, Y., Kong, C., Yu, L., Li, P., Chen, C., Zheng, X., Hertzberg, V.S., Yang, C.: 4sdrug: Symptom-based set-to-set small and safe drug recommendation. In: Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, pp. 3970–3980 (2022)
- [186] Wu, Z., Yao, H., Su, Z., Liebovitz, D.M., Glass, L.M., Zou, J., Finn, C., Sun, J.: Knowledge-driven new drug recommendation. *arXiv preprint arXiv:2210.05572* (2022)
- [187] Gao, Z., Gao, M., Chen, C.-h., Zhou, Y., Zhan, Z.-H., Ren, Y.: Knowledge graph of wastewater-based epidemiology development: A data-driven analysis based on research topics and trends. *Environmental Science and Pollution Research* **30**(11), 28373–28382 (2023)
- [188] Domingo-Fernández, D., Baksi, S., Schultz, B., Gadiya, Y., Karki, R., Raschka, T., Ebeling, C., Hofmann-Apitius, M., Kodamullil, A.T.: Covid-19 knowledge

- graph: a computable, multi-modal, cause-and-effect knowledge model of covid-19 pathophysiology. *Bioinformatics* **37**(9), 1332–1334 (2021)
- [189] Turki, H., Jemielniak, D., Taieb, M.A.H., Gayo, J.E.L., Aouicha, M.B., Banat, M., Shafee, T., Prud’hommeaux, E., Lubiana, T., Das, D., *et al.*: Using logical constraints to validate statistical information about disease outbreaks in collaborative knowledge graphs: the case of covid-19 epidemiology in wikidata. *PeerJ Computer Science* **8**, 1085 (2022)
 - [190] Pressat Laffouilhère, T., Grosjean, J., Pinson, J., Darmoni, S.J., Leveque, E., Lanoy, E., Bénichou, J., Soualmia, L.F.: Ontological representation of causal relations for a deep understanding of associations between variables in epidemiology. In: *Artificial Intelligence in Medicine: 20th International Conference on Artificial Intelligence in Medicine, AIME 2022, Halifax, NS, Canada, June 14–17, 2022, Proceedings*, pp. 47–56 (2022). Springer
 - [191] Fecho, K., Bizon, C., Miller, F., Schurman, S., Schmitt, C., Xue, W., Morton, K., Wang, P., Tropsha, A., *et al.*: A biomedical knowledge graph system to propose mechanistic hypotheses for real-world environmental health observations: cohort study and informatics application. *JMIR Medical Informatics* **9**(7), 26714 (2021)
 - [192] Wolffe, T.A., Vidler, J., Halsall, C., Hunt, N., Whaley, P.: A survey of systematic evidence mapping practice and the case for knowledge graphs in environmental health and toxicology. *Toxicological Sciences* **175**(1), 35–49 (2020)
 - [193] Wu, J.: Construct a knowledge graph for china coronavirus (covid-19) patient information tracking. *Risk Management and Healthcare Policy*, 4321–4337 (2021)
 - [194] Yu, G., Tabatabaei, M., Mezei, J., Zhong, Q., Chen, S., Li, Z., Li, J., Shu, L., Shu, Q.: Improving chronic disease management for children with knowledge graphs and artificial intelligence. *Expert Systems with Applications* **201**, 117026 (2022)
 - [195] Cao, L., Zhang, H., Feng, L.: Building and using personal knowledge graph to improve suicidal ideation detection on social media. *IEEE Transactions on Multimedia* **24**, 87–102 (2020)
 - [196] Liu, H., Li, Y., Hong, R., Li, Z., Li, M., Pan, W., Glowacz, A., He, H.: Knowledge graph analysis and visualization of research trends on driver behavior. *Journal of Intelligent & Fuzzy Systems* **38**(1), 495–511 (2020)
 - [197] Wang, S., Zhang, Y., Lin, B., Li, B.: Interpretable emotion analysis based on knowledge graph and occ model. In: *Proceedings of the 31st ACM International Conference on Information & Knowledge Management*, pp. 2038–2045 (2022)
 - [198] Singhal, K., Azizi, S., Tu, T., Mahdavi, S.S., Wei, J., Chung, H.W., Scales, N.,

- Tanwani, A., Cole-Lewis, H., Pfohl, S., et al.: Large language models encode clinical knowledge. arXiv preprint arXiv:2212.13138 (2022)
- [199] Nath, S., Marie, A., Ellershaw, S., Korot, E., Keane, P.A.: New meaning for nlp: the trials and tribulations of natural language processing with gpt-3 in ophthalmology. *British Journal of Ophthalmology* **106**(7), 889–892 (2022)
- [200] Moor, M., Banerjee, O., Abad, Z.S.H., Krumholz, H.M., Leskovec, J., Topol, E.J., Rajpurkar, P.: Foundation models for generalist medical artificial intelligence. *Nature* **616**(7956), 259–265 (2023)